

# Exercises in Parallel Programming in Fortran and C/C++

Dr. Sara Collins

SoSe 2016

## Sheet 6

---

### Exercise 1: Integration routine with OpenMP (3 Points)

In exercise 4 from sheet 5 you modified your integration module to include the OpenMP directive `omp parallel`.


- (a) Produce a new version of the integration routine using the OpenMP directive `omp do`.
- (b) Test that the routine produces the correct results and compare the wall clock times for the previous version (exercise 4 sheet 5) and the new one (using `omp do`).

### Exercise 2: Monte-Carlo Integration (8 Points)

So far we have only considered 1-dimensional integrals. However, scientific problems usually require numerical computation of  $n$ -dimensional integrals, where  $n$  is in the range of  $n = 2$  (in the simplest case) to  $n = \text{very very large}$ . Especially for the latter, the classical methods of integration (trapezoidal rule, rectangle rule, ...) are unsuitable and one uses so-called *Monte Carlo methods*. These methods use random numbers to estimate the (high-dimensional) integral. The precision of the result increases with the number of random numbers.

We first consider a simple example — the calculation of a quadrant of the area of a circle of radius  $R = 1.0$  (also known as Monte Carlo calculation of the number  $\pi$ ). One generates  $N$  (eg  $N = 10000$ ) pairs  $r_i = (x_i, y_i)$  (with  $i = 1, \dots, N$ ) of uniformly distributed random numbers  $x_i$  and  $y_i$  where  $0 < x_i, y_i \leq 1$ . One counts how many pairs are inside the quadrant of the circle ( $R^2 \geq x_i^2 + y_i^2$ ), denoted  $N_{in}$ , and uses

$$\frac{N_{in}}{N} \approx \frac{A_{\text{Circle}}/4}{A_{\square}} \quad (1)$$



where  $A_{\square} = 1.0$ .

**Hint:** Use the formula  $A_{\text{Circle}} = \pi R^2$  in order to estimate  $\pi$  using this method.

- (a) Write a program that calculates the area of a circle ( $n = 2$ ) with the above described Monte-Carlo method. Display in a plot how the estimate of  $\pi$  approaches the exact result as  $N$  is increased.
- (b) Generalise your program so that it works for  $n$  dimensions – calculating the volume of an  $n$ -dimensional sphere ( $A_V$ ). Equation (2) is modified to

$$\frac{N_{in}}{N} \approx \frac{A_V/2^n}{A_{\square}} \quad (2)$$

Compare it with the exact result given by:

$$V = \begin{cases} \frac{\pi^k}{k!} & \text{für } n = 2k \\ \frac{2^{k+1}\pi^k}{(2k+1)!!} & \text{für } n = 2k + 1 \end{cases} \quad (3)$$

wobei  $(2k + 1)!! = 1 \cdot 3 \cdots (2k - 1) \cdot (2k + 1)$

- (c) Parallelise your program with OpenMP and test that it provides the correct functionality.

### Exercise 3: sparse matrix vector multiplication (10 Points)

Many scientific problems involve matrix-vector multiplication:  $\vec{w} = M\vec{v}$ . We consider here the case of a matrix of size  $n \times n$ . If the matrix is sparse (i.e. mostly contains zeros) then the standard representation of the matrix (as a two-dimensional array) and the implementation of the matrix-vector multiplication (with two nested `do` loops) is very inefficient.

One possibility is to represent the matrix in compressed form using a vector which only contains the non-zero entries using the Compressed Row Storage (CRS) format. This requires three vectors - one holding the non-zero entries (`val`), the second storing the column indices of the non-zero entries (`col`) and the third (`row`) storing the indices of the elements in `val` which start a new row. For example,

$$\begin{pmatrix} 10 & 0 & 0 & 0 & -2 & 0 \\ 3 & 9 & 0 & 0 & 0 & 3 \\ 0 & 7 & 8 & 7 & 0 & 0 \\ 3 & 0 & 8 & 7 & 5 & 0 \\ 0 & 8 & 0 & 9 & 9 & 13 \\ 0 & 4 & 0 & 0 & 2 & -1 \end{pmatrix} \quad (4)$$

$$\text{val} = (10, -2, 3, 9, 3, 7, 8, 7, 3, 8, 7, 5, 8, 9, 9, 13, 4, 2, -1) \quad (5)$$

$$\text{col} = (1, 5, 1, 2, 6, 2, 3, 4, 1, 3, 4, 5, 2, 4, 5, 6, 2, 5, 6) \quad (6)$$

$$\text{row} = (1, 3, 6, 9, 13, 17, 20) \quad (7)$$

If there are  $n_{\text{zero}}$  non-zero entries in  $M$  ( $n_{\text{zero}} = 19$  in the example), then `val` and `col` are of length  $n_{\text{zero}}$  and `row` is of length  $n+1$  (7 in the example), `row(n+1) = nzero+1`. The matrix-vector multiplication can be performed using

```
do i = 1, n
  do j = row(i), row(i+1)-1
    w(i) = w(i) + val(j) * v(col(j))
  enddo
enddo
```

- (a) Write a program which performs matrix-vector multiplication using the standard implementation (using a two-dimensional array for  $M$ ). Extend the program to convert  $M$  into CRS format and perform the matrix-vector multiplication again. Check that the correct functionality is reproduced. You should initialise  $M$  to be a sparse matrix, for example, only containing 1% non-zero entries ( $n_{\text{zero}} \approx n^2/100$ ).
- (b) Parallelize your program for both cases using OpenMP. Your program should be able to compile with and without OpenMP.
- (c) Measure the wallclock time for the multiplication (standard and CRS) for different values of  $n$  (from small to very large keeping  $n_{\text{zero}}/n^2 \approx \text{constant}$ ), with and without OpenMP and using a different number of threads. Display the results graphically as a function of  $n$ . Give also the performance (in Flops/s) and the efficiency as a function of  $n$ .
- (d) Test different scheduling options for the do loops.